

## Libraries

In [1]:

```
# Install pycocotools
#!pip install pycocotools

import os
import json
import numpy as np
import tensorflow as tf
from tensorflow.keras import layers, models, losses
from tensorflow.keras.optimizers import Adam
from pycocotools.coco import COCO
from sklearn.model_selection import train_test_split
import cv2
import matplotlib.pyplot as plt
from google.colab import drive
from PIL import Image
import tensorflow as tf
from tensorflow.keras.metrics import BinaryAccuracy, MeanSquaredError
import random
from tensorflow.keras.metrics import Metric

from tensorflow.keras import Model
from tensorflow.keras.layers import (
    Input, Conv2D, BatchNormalization, LeakyReLU, MaxPooling2D, Flatten, Dense, GlobalAveragePooling2D
)

from tensorflow.keras import mixed_precision
mixed_precision.set_global_policy('mixed_float16') # all operations during training will
be done with this precision (matrix mul...)
# speeds up training time and inference time
```

In [24]:

```
names_of_class= ["FLY", "NOPALM", "PAUSE", "UNDEFINED"]

def coco_to_center(coco_bbox):
    """
    Convert COCO format bbox [x, y, width, height] to center-based format [center_x, center_y, width, height].

    Args:
        coco_bbox (list or tuple): Bounding box in COCO format [x, y, width, height].

    Returns:
        list: Bounding box in center-based format [center_x, center_y, width, height].
    """
    x, y, width, height = coco_bbox
    center_x = x + width / 2
    center_y = y + height / 2
    return [center_x, center_y, width, height]

def center_to_coco(center_bbox):
    """
    Convert center-based format bbox [center_x, center_y, width, height] to COCO format [x, y, width, height].

    Args:
        center_bbox (list or tuple): Bounding box in center-based format [center_x, center_y, width, height].

    Returns:
        list: Bounding box in COCO format [x, y, width, height].
    """
    center_x, center_y, width, height = center_bbox
```

```

x = center_x - width / 2
y = center_y - height / 2
return [x, y, width, height]

def plot_image_with_bbox(image, true_bbox, true_class, pred_bbox, pred_class, bbox_format
="coco"):

    plt.figure(figsize=(6, 6))

    # Multiply by 255 to rescale the image to [0, 255] range
    image = image * 255 # Undo normalization

    # Ensure image is in the correct dtype (uint8) for display
    image = image.astype('uint8')

    plt.imshow(image, cmap='gray')
    plt.axis('off')

    # Convert bounding boxes based on the format
    if bbox_format == "center":
        true_bbox = center_to_coco(true_bbox)
        true_bbox = np.array(true_bbox, dtype=np.float16)
        pred_bbox = center_to_coco(pred_bbox)
        pred_bbox = np.array(pred_bbox, dtype=np.float16)

    # Draw the true bounding box if it exists
    if not (true_bbox == [0, 0, 0, 0]).all():
        x, y, w, h = true_bbox
        plt.gca().add_patch(plt.Rectangle((x, y), w, h, linewidth=2, edgecolor='r', face
color='none'))
        #plt.gca().add_patch(plt.Rectangle((x * 360, y * 240), w * 360, h * 240, linewidth=2, edgecolor='r', facecolor='none'))

    # Draw the predicted bounding box if it exists
    if not (pred_bbox == [0, 0, 0, 0]).all():
        x, y, w, h = pred_bbox
        plt.gca().add_patch(plt.Rectangle((x, y), w, h, linewidth=2, edgecolor='g', face
color='none'))

    # Add title with both true and predicted class
    plt.title(f"True Class: {true_class} [{names_of_class[np.argmax(true_class)}], Predi
cted Class: {pred_class} [{names_of_class[np.argmax(pred_class)}]")
    print("IoU:", 1 - iou_loss(true_bbox, pred_bbox, "coco"))
    plt.show()

```

In [3]:

```

# Mount Google Drive
drive.mount('/content/drive')

```

Mounted at /content/drive

In [4]:

```

# HYPERPARAMETERS

include_MIRROR = True
# - double the dataset by adding mirror images

bbox_format = "center" # "coco" - top left, "center" - center
# - do you want to use default COCO bbox format, or switch to Center bbox format that cou
ld potentially improve object detecting

include_MOVM = False

```

In [5]:

```

# Ova skripta učitava podatke i pretvara ih u format pogodan za treniranje

# Path to your COCO dataset files
annotations_path = "/content/drive/MyDrive/Colab Notebooks/FINAL_PALM/result.json"

```

```

images_path = "/content/drive/MyDrive/Colab Notebooks/FINAL_PALM"
if not os.path.exists(images_path):
    print(f"Image {images_path} does not exist!")

# Load COCO dataset
coco = COCO(annotations_path)

# Extract categories
categories = coco.loadCats(coco.getCatIds())
category_names = [cat['name'] for cat in categories]
print("Categories:", category_names)
print()

# Load image and annotation data
image_ids = coco.getImgIds()
images = coco.loadImgs(image_ids)

# Prepare data for Palm and No Palm
data = []
for image_info in images: #[:int(len(images)/2)]:
    img_id = image_info['id']
    img_path = os.path.join(images_path, image_info['file_name'].replace("\\", "/"))
    movm_path2 = "movement\\" + image_info['file_name'].replace("images", "movement").split("-", 1)[1].replace(".png", "_movement.png")
    movm_path = os.path.join(images_path, movm_path2.replace("\\", "/"))

    # print("Image file name:", image_info['file_name'])
    # print("MOVM file name:", movm_path2)
    # print(f"Image path: {img_path}")
    # print(f"MOVM path: {movm_path}")
    # print()
    anns = coco.loadAnns(coco.getAnnIds(imgIds=[img_id]))

    # Check if Palm exists
    palm_annotations = [ann for ann in anns if (ann['category_id'] != 1)]
    if palm_annotations:
        for ann in palm_annotations:
            bbox = ann['bbox'] # [x, y, width, height]
            if bbox_format == "center":
                bbox = coco_to_center(bbox)
            width, height = Image.open(img_path).size

            # Add original data
            data.append({
                "image_path": img_path,
                "movement_path": movm_path,
                "label": ann['category_id'], # FLY=0 or PAUSE=2 or UNDEFINED=3
                "bbox": bbox
            })
    else:
        # Add original "No Palm" data
        data.append({
            "image_path": img_path,
            "movement_path": movm_path,
            "label": 1, # NOPALM=1
            "bbox": [0, 0, 0, 0]
        })

def preprocess_image(image_path, movm_path, bbox, img_size=(240, 360), flip_or_not=False):
    # (480, 720)

    # Load image in grayscale
    img = cv2.imread(image_path, cv2.IMREAD_GRAYSCALE)
    movm_img = None

    if movm_path is not None:
        movm_img = cv2.imread(movm_path, cv2.IMREAD_GRAYSCALE)

    if movm_img is None:
        print(f"Warning: Image at {movm_path} could not be loaded.")

```

```

        return None, None, None

    # Resize and normalize the image to [0, 1]
    movm_img = cv2.resize(movm_img, (img_size[1], img_size[0])) / 255.0

    if flip_or_not:
        movm_img = cv2.flip(movm_img, 1)

    if img is None:
        print(f"Warning: Image at {image_path} could not be loaded.")
        return None, None, None

    # Resize and normalize the image to [0, 1]
    img = cv2.resize(img, (img_size[1], img_size[0])) / 255.0 # Note: width comes first
    in OpenCV resize

    # If Palm, normalize bounding box
    if bbox:
        h, w = img_size
        x, y, bw, bh = bbox
        #bbox = [x / w, y / h, bw / w, bh / h] # Normalize to [0, 1]
    else:
        bbox = [0, 0, 0, 0] # No bounding box for "No Palm"

    # Apply horizontal flip if requested
    if flip_or_not:
        img = cv2.flip(img, 1) # Horizontal flip

    # Flip bounding box coordinates
    if bbox != [0, 0, 0, 0]: # Only flip if there's a valid bounding box
        if bbox_format == "coco":
            #FLIP FOR COCO
            bbox = [
                img_size[1] - (bbox[0] + bbox[2]), # Flip x by mirroring x + width
                bbox[1], # y remains the same
                bbox[2], # Width remains the same
                bbox[3] # Height remains the same
            ]
        elif bbox_format == "center":
            # FLIP FOR CENTER BASED BBOX
            bbox = [
                img_size[1] - bbox[0], # Flip x by mirroring x + width
                bbox[1], # y remains the same
                bbox[2], # Width remains the same
                bbox[3] # Height remains the same
            ]
    # Cast bbox and label to float16
    bbox = np.array(bbox, dtype=np.float16)

    return img, movm_img, bbox

```

loading annotations into memory...  
 Done (t=0.94s)  
 creating index...  
 index created!  
 Categories: ['FLY', 'NOPALM', 'PAUSE', 'UNDEFINED']

In [26]:

```

# Učitavanje podataka, dijeljenje u train/test data, te na kraju definiranje arhitekture
NN. (Nikako ne koristiti batchnormalization layer)

def data_generator(data, img_size=(240, 360), is_include_MOVM=False):
    images, movements, labels, bboxes = [], [], [], []
    for entry in data:
        img, movm_img, bbox = preprocess_image(entry['image_path'], entry['movement_path']

```

```

'] if include_MOVM else None, entry['bbox'], img_size, flip_or_not = False)
    if img is not None:
        images.append(img)
        movements.append(movm_img)
        labels.append(entry['label'])
        bboxes.append(bbox)

    if include_MIRROR:
        f_img, f_movm_img, f_bbox = preprocess_image(entry['image_path'], entry['movement_path'] if include_MOVM else None, entry['bbox'], img_size, flip_or_not = True)
        if img is not None:
            images.append(f_img)
            movements.append(f_movm_img)
            labels.append(entry['label'])
            bboxes.append(f_bbox)

# Convert to numpy arrays and add channel dimension
images = np.array(images, dtype=np.float16)[..., np.newaxis]

if is_include_MOVM:
    movements = np.array(movements, dtype=np.float16)[..., np.newaxis]
    images = np.concatenate([images, movements], axis=-1)

# Combine label and bbox for the output (first element is label, next 4 are bbox)
outputs = []
for label, bbox in zip(labels, bboxes):
    # Ensure bbox is a 4-element vector (x, y, width, height)
    assert len(bbox) == 4, f"Expected bbox to have 4 elements, got {len(bbox)}"

    # One-hot encode the label (for 4 classes)
    one_hot_label = np.eye(4)[label]

    # Combine the one-hot label with the bounding box
    output = np.concatenate([one_hot_label, bbox]).astype(np.float16)
    outputs.append(output)

# Convert outputs to numpy array with dtype float16
return images, np.array(outputs, dtype=np.float16)

#####

# Prepare training and testing data
IMG_SIZE = (240, 360)
X_all, y_all = data_generator(data, IMG_SIZE, is_include_MOVM = include_MOVM)
#split train/test 70-30
X_train, X_test, y_train, y_test = train_test_split(X_all, y_all, test_size=0.3, random_state=42, shuffle=True)

# Extract labels from y_all
labels = y_all[:, 0] # Assuming the first element is the label

# Count the occurrences of each label
unique_labels, label_counts = np.unique(labels, return_counts=True)

# Print the results
for label, count in zip(unique_labels, label_counts):
    print(f"Label {label}: {count} samples")

print()
print(f"Training: {len(X_train)} instances, {len(y_train)} labels | Testing: {len(X_test)} instances, {len(y_test)} labels")
random_indices = random.sample(range(len(y_train)), 10)

# Print y_train and bbox_train for the selected instances
for idx in random_indices:
    print(f"Instance {idx}: Label = {y_train[idx]}")

print("\nX_train shape:", X_train.shape) # Should be (N, 480, 720, 1)
print("y_train shape:", y_train.shape) # Should be (N, 5), where 5 is [label, bbox_x, bb

```

```

ox_y, bbox_w, bbox_h]
print("X_test shape:", X_test.shape) # Should be (N, 480, 720, 1)
print("y_test shape:", y_test.shape) # Should be (N, 5), where 5 is [label, bbox_x, bbox
_y, bbox_w, bbox_h]

# Define the combined model
# def build_model(input_shape):
#     input_img = layers.Input(shape=input_shape)

#     # Feature extraction
#     # x = layers.Conv2D(8, (5, 5), activation="relu")(input_img)
#     # x = layers.MaxPooling2D((6, 6))(x)
#     # x = layers.Conv2D(24, (3, 3), activation="relu")(x)
#     # x = layers.MaxPooling2D((2, 2))(x)
#     # x = layers.Flatten()(x)
#     # x = layers.Dense(8, activation="relu")(x)

#     x = layers.Conv2D(2, (4, 4), activation="relu")(input_img)
#     x = layers.MaxPooling2D((2, 2))(x)
#     x = layers.Conv2D(4, (3, 3), activation="relu")(input_img)
#     x = layers.MaxPooling2D((6, 6))(x)
#     x = layers.Conv2D(16, (3, 3), activation="relu")(x)
#     x = layers.Flatten()(x)
#     x = layers.Dense(8, activation="relu")(x)

#     # Single output layer for both classification and bounding box
#     output = layers.Dense(5, activation="sigmoid", name="output")(x)

#     model = models.Model(inputs=input_img, outputs=output)
#     model.summary()

#     return model

#CUSTOM NN
#def build_model(input_shape):
#    input_img = layers.Input(shape=input_shape)

#    # x = layers.Conv2D(32, (3, 3), activation='relu', padding='same')(input_img)
#    # x = layers.MaxPooling2D((2, 2))(x)

#    # # 2nd Conv Block
#    # x = layers.Conv2D(32, (2, 2), activation='relu', padding='same')(x)
#    # x = layers.MaxPooling2D((2, 2))(x)

#    # # 3rd Conv Block
#    # x = layers.Conv2D(64, (3, 3), activation='relu', padding='same')(x)
#    # x = layers.MaxPooling2D((2, 2))(x)

#    # # 4th Conv Block
#    # x = layers.Conv2D(128, (2, 2), activation='relu', padding='same')(x)
#    # x = layers.MaxPooling2D((3, 3))(x)

#    # # Flatten the output for fully connected layers
#    # x = layers.Flatten()(x)

#    # # Fully Connected Layer
#    # x = layers.Dense(32, activation='relu')(x)
#    # x = layers.Dropout(0.15)(x) # Dropout to prevent overfitting

#####
# YOLO V1 NN
def build_model(input_shape, grid_size=7, num_classes=4, isMOVM = False):

    # without MOVm images
    if isMOVM == False:

        input_img = Input(shape=input_shape)

        # Convolutional Backbone
        x = Conv2D(64, (7, 7), strides=2, padding='same', activation=None)(input_img)

```

```

#x = BatchNormalization()(x)
x = LeakyReLU(alpha=0.1)(x)
x = MaxPooling2D(pool_size=(2, 2), strides=2)(x)

x = Conv2D(128, (3, 3), strides=1, padding='same', activation=None)(x)
#x = BatchNormalization()(x)
x = LeakyReLU(alpha=0.1)(x)
x = MaxPooling2D(pool_size=(2, 2), strides=2)(x)

x = Conv2D(64, (1, 1), strides=1, padding='same')(x)
x = LeakyReLU(alpha=0.1)(x)
x = Conv2D(128, (3, 3), strides=1, padding='same')(x)
x = LeakyReLU(alpha=0.1)(x)
x = Conv2D(128, (1, 1), strides=1, padding='same')(x)
x = LeakyReLU(alpha=0.1)(x)
x = Conv2D(128, (3, 3), strides=1, padding='same')(x)
#x = BatchNormalization()(x)
x = LeakyReLU(alpha=0.1)(x)
x = MaxPooling2D(pool_size=(2, 2), strides=2)(x)

# Concatination of convolutional info
x = Flatten()(x)
#x = GlobalAveragePooling2D()(x) ILI PROBAT global max pool!!!!

# This Fully connected layer will be used only by bbox pred head
x = Dense(128, activation='relu')(x)

# This Fully connected layer will be used only by classification head
x_class = Dense(64, activation='relu')(x)
#x_class = Dense(32, activation='relu')(x_class)

# Bounding box predictions
bbox_output = Dense(4, activation='relu', name='bbox_output')(x)

# Class probabilities
class_output = Dense(4, activation='softmax', name='class_output')(x_class)

# Concatenate bbox and class probabilities into a single 8-element output
output = tf.keras.layers.Concatenate(name='output')([class_output, bbox_output])

model = models.Model(inputs=input_img, outputs=output)
# model = Model(inputs=input_img, outputs=[class_output, combined_output])

#with MOVIM images
else:
    input_img1 = Input(shape=input_shape, name="input_image1") # Black-and-white im
age
    input_img2 = Input(shape=input_shape, name="input_image2") # Difference image

    # Branch for input image 1
    x1 = Conv2D(64, (7, 7), strides=2, padding='same', activation=None)(input_img1)
    x1 = LeakyReLU(alpha=0.1)(x1)
    x1 = MaxPooling2D(pool_size=(2, 2), strides=2)(x1)

    x1 = Conv2D(128, (3, 3), strides=1, padding='same', activation=None)(x1)
    x1 = LeakyReLU(alpha=0.1)(x1)
    x1 = MaxPooling2D(pool_size=(2, 2), strides=2)(x1)

    # Branch for input image 2
    x2 = Conv2D(64, (7, 7), strides=2, padding='same', activation=None)(input_img2)
    x2 = LeakyReLU(alpha=0.1)(x2)
    x2 = MaxPooling2D(pool_size=(2, 2), strides=2)(x2)

    x2 = Conv2D(128, (3, 3), strides=1, padding='same', activation=None)(x2)
    x2 = LeakyReLU(alpha=0.1)(x2)
    x2 = MaxPooling2D(pool_size=(2, 2), strides=2)(x2)

    # Merge the features from both branches
    merged = tf.keras.layers.Concatenate()([x1, x2])

    x = Conv2D(64, (1, 1), strides=1, padding='same')(merged)

```

```

x = LeakyReLU(alpha=0.1)(x)
x = Conv2D(128, (3, 3), strides=1, padding='same')(x)
x = LeakyReLU(alpha=0.1)(x)
x = Conv2D(128, (1, 1), strides=1, padding='same')(x)
x = LeakyReLU(alpha=0.1)(x)
x = Conv2D(128, (3, 3), strides=1, padding='same')(x)
#x = BatchNormalization()(x)
x = LeakyReLU(alpha=0.1)(x)
x = MaxPooling2D(pool_size=(2, 2), strides=2)(x)

# Concatination of convolutional info
x = Flatten()(x)

# This Fully connected layer will be used only by bbox pred head
x_local = Dense(128, activation='relu')(x)

# This Fully connected layer will be used only by classification head
x_class = Dense(64, activation='relu')(x)
#x_class = Dense(32, activation='relu')(x_class)

# Class probabilities
class_output = Dense(num_classes, activation='softmax', name='class_output')(x_c
lass)

# Bounding box predictions
bbox_output = Dense(4, activation='relu', name='bbox_output')(x_local)

# Concatenate bbox and class probabilities into a single 8-element output
output = tf.keras.layers.Concatenate(name='output')([class_output, bbox_output])

model = models.Model(inputs=[input_img1, input_img2], outputs=output)

model.summary()
return model

```

Label 0.0: 696 samples  
Label 1.0: 742 samples

Training: 1006 instances, 1006 labels | Testing: 432 instances, 432 labels

Instance 154:	Label = [ 1. 0. 0. 0. 98.5 63.44 31.44 52.53]
Instance 678:	Label = [ 1. 0. 0. 0. 91.56 134.4 88.7 68.4 ]
Instance 696:	Label = [ 1. 0. 0. 0. 250.2 147.6 47.78 22.55]
Instance 172:	Label = [ 0. 0. 1. 0. 187.1 69.7 37.47 57.06]
Instance 313:	Label = [ 0. 0. 1. 0. 148.4 102. 90.5 104.25]
Instance 775:	Label = [ 1. 0. 0. 0. 73.3 120.75 49.16 42.72]
Instance 693:	Label = [ 1. 0. 0. 0. 301.8 137.2 106.75 42.34]
Instance 74:	Label = [ 0. 0. 0. 1. 109.6 99.8 30.44 45.94]
Instance 434:	Label = [ 0. 0. 0. 1. 255.8 181.1 48.3 36.4]
Instance 33:	Label = [ 1. 0. 0. 0. 247.6 160.9 51.38 31.05]

X\_train shape: (1006, 240, 360, 1)  
y\_train shape: (1006, 8)  
X\_test shape: (432, 240, 360, 1)  
y\_test shape: (432, 8)

In [15]:

```

# Ovo je najvažniji dio, odnosi se na treniranje neuronske mreže, definiranje njene optim
izacije (loss funct.) i metrika da pratimo napredak kroz treniranje
# Nebi trebalo ništa mijenjati, eventualno epochs (koliko rundi treniranja) i learning ra
te (lr), ili izmjeniti samu arhitektutu NN u prethodnim dijelovima
# Loss function je kombinacija mse * iou
# TIP: ako u prvih 10-15 epocha IoU metric bude i dalje 0.0, onda prekini izvođenje koda
i pokreni iznova, dok se ne pojavi bar neki broj u prvih 10 epocha

from tensorflow.keras.callbacks import ReduceLROnPlateau, EarlyStopping
#tf.config.run_functions_eagerly(True)
#tf.data.experimental.enable_debug_mode()

classification_loss_fn = tf.keras.losses.SparseCategoricalCrossentropy()

```



```

# num_classes
def iou_loss(y_true, y_pred, bbox_format_rn=None):
    """
    Compute IoU Loss between true and predicted bounding boxes.
    y_true and y_pred are expected to be in the format [x, y, w, h].
    """
    if bbox_format_rn is None:
        bbox_format_rn=bbox_format

    y_true = tf.cast(y_true, tf.float32)
    y_pred = tf.cast(y_pred, tf.float32)

    if bbox_format_rn == "center":
        # Convert center-based bbox [cx, cy, w, h] to corner-based bbox [x_min, y_min, x_max, y_max]
        y_true_x_min = y_true[..., 0] - y_true[..., 2] / 2
        y_true_y_min = y_true[..., 1] - y_true[..., 3] / 2
        y_true_x_max = y_true[..., 0] + y_true[..., 2] / 2
        y_true_y_max = y_true[..., 1] + y_true[..., 3] / 2

        y_pred_x_min = y_pred[..., 0] - y_pred[..., 2] / 2
        y_pred_y_min = y_pred[..., 1] - y_pred[..., 3] / 2
        y_pred_x_max = y_pred[..., 0] + y_pred[..., 2] / 2
        y_pred_y_max = y_pred[..., 1] + y_pred[..., 3] / 2

    elif bbox_format_rn == "coco":
        # For COCO format, use the bbox directly
        y_true_x_min = y_true[..., 0]
        y_true_y_min = y_true[..., 1]
        y_true_x_max = y_true[..., 0] + y_true[..., 2]
        y_true_y_max = y_true[..., 1] + y_true[..., 3]

        y_pred_x_min = y_pred[..., 0]
        y_pred_y_min = y_pred[..., 1]
        y_pred_x_max = y_pred[..., 0] + y_pred[..., 2]
        y_pred_y_max = y_pred[..., 1] + y_pred[..., 3]
    else:
        raise ValueError("Invalid bbox_format. Use 'coco' or 'center'.")

    # Intersection coordinates
    inter_x_min = tf.maximum(y_true_x_min, y_pred_x_min)
    inter_y_min = tf.maximum(y_true_y_min, y_pred_y_min)
    inter_x_max = tf.minimum(y_true_x_max, y_pred_x_max)
    inter_y_max = tf.minimum(y_true_y_max, y_pred_y_max)

    # Intersection area
    inter_area = tf.maximum(0.0, inter_x_max - inter_x_min) * tf.maximum(0.0, inter_y_max - inter_y_min)

    # Union area
    true_area = (y_true_x_max - y_true_x_min) * (y_true_y_max - y_true_y_min)
    pred_area = (y_pred_x_max - y_pred_x_min) * (y_pred_y_max - y_pred_y_min)
    union_area = true_area + pred_area - inter_area

    # IoU
    iou = inter_area / tf.maximum(union_area, 1e-10) # Avoid division by zero

    # IoU Loss
    return 1 - iou

def total_loss(y_true, y_pred):

    # print("y_true shape:", y_true.shape)
    # print("y_pred shape:", y_pred.shape)

    ### Print the first few elements of y_true and y_pred to inspect
    # tf.print("y_true values:", y_true) # Only the first 5 elements
    # tf.print("y_pred values:", y_pred) # Only the first 5 elements

    class_true = y_true[..., 0:4]

```

```

bbox_true = y_true[..., 4:]

class_pred = y_pred[..., 0:4]
bbox_pred = y_pred[..., 4:]

# Create a mask to separate Palm and No Palm
class_mask = tf.cast(tf.logical_or( tf.argmax(class_true, axis=-1) == 0, tf.argmax(class_true, axis=-1) == 2), tf.float32)

squared_errors = tf.square(bbox_true - bbox_pred) # Shape: (batch_size, 4)

# Per-sample loss (mean across bbox coordinates)
per_sample_loss = tf.reduce_mean(squared_errors, axis=-1) # Shape: (batch_size,)

# Batch-level loss (mean across the batch)
batch_loss = tf.reduce_mean(per_sample_loss)

total_bbox_loss = batch_loss

# For No Palm (label=0), only calculate class loss
# For Palm (label=1), calculate both class and bbox loss
abs_bbox_loss = total_bbox_loss * class_mask # Apply mask to bbox loss (ignore bbox loss if NOPALM or UNDEFINED)

class_true_labels = tf.argmax(class_true, axis=-1)
classification_loss = classification_loss_fn(class_true_labels, class_pred)
# classification_loss = tf.equal(class_true, class_pred)
# classification_loss = tf.cast(classification_loss, tf.float32)

# tf.print("class_true:", class_true)
# tf.print("class_pred:", class_pred)
# tf.print("classification_loss:", classification_loss, "\n")

# IoU loss
bbox_iou_loss = iou_loss(bbox_true, bbox_pred, bbox_format)
bbox_iou_loss = bbox_iou_loss * class_mask

# Total loss
#return 0 * classification_loss + 0 * bbox_iou_loss * 100 + 1 * abs_bbox_loss
#return 0 * classification_loss + bbox_iou_loss * abs_bbox_loss * 10/7 #OVA JE GOAT
return 2 * classification_loss**2 + 1 * (bbox_iou_loss+1)**2 * abs_bbox_loss #1000 * classification_loss**4 +

# Define the custom Sparse Categorical Crossentropy loss function
def custom_scc_loss(y_true, y_pred):
    """
    Custom Sparse Categorical Crossentropy Loss.
    y_true: One-hot encoded true labels (batch_size, num_classes).
    y_pred: Predicted probabilities (batch_size, num_classes).
    """

    class_true = y_true[..., 0:4]
    #bbox_true = y_true[..., 4:]

    class_pred = y_pred[..., 0:4]
    #bbox_pred = y_pred[..., 4:]

    # Compute SCCE
    # class_true_labels = tf.argmax(class_true, axis=-1)
    # loss = classification_loss_fn(class_true_labels, class_pred)
    loss = -tf.reduce_sum(class_true * tf.math.log(class_pred + 1e-7), axis=-1) # Add epsilon for numerical stability
    return tf.reduce_mean(loss)

def iou_metric(y_true, y_pred):
    class_true = y_true[..., 0:4]
    bbox_true = y_true[..., 4:]
    bbox_pred = y_pred[..., 4:]

```

```

iou = 1 - iou_loss(bbox_true, bbox_pred) # Since `iou_loss` returns 1 - IoU

# Apply the condition: if true class is 1 or 3, set IoU to 0
condition = tf.logical_or(tf.argmax(class_true, axis=-1) == 1, tf.argmax(class_true,
axis=-1) == 3)
iou = tf.where(condition, 0.0, iou)

return iou * (10/7) # -> *(10/7) zato jer 30% podataka class(NOPALM ili UNDEFINED) p
a kvare statistiku za cca ovaj omjer

def class_accuracy(y_true, y_pred):

    class_true = y_true[..., 0:4]
    class_pred = y_pred[..., 0:4]

    # Compare the true class with the predicted class
    class_true_index = tf.cast(tf.argmax(class_true, axis=-1), tf.float32)
    class_pred_index = tf.cast(tf.argmax(class_pred, axis=-1), tf.float32)

    # Compare the predicted class index with the true class index
    correct_predictions = tf.equal(class_true_index, class_pred_index)

    # Calculate accuracy: the number of correct predictions divided by the total number o
f samples
    accuracy = tf.cast(correct_predictions, tf.float32)
    #accuracy = tf.reduce_mean(tf.cast(correct_predictions, tf.float32))

    return accuracy

def mse_metric(y_true, y_pred):
    """
    Custom MSE metric with conditional logic.
    If y_true[..., 0] is 1 or 3, the MSE for those samples is set to 0.
    """
    # Extract bounding box coordinates and cast to float32
    bbox_true = tf.cast(y_true[..., 4:], tf.float32)
    bbox_pred = tf.cast(y_pred[..., 4:], tf.float32)
    # bbox_true = tf.cast(y_true[..., 4:], tf.float32)
    # bbox_pred = tf.cast(y_pred[..., 4:], tf.float32)
    class_true = y_true[..., 0:4]

    # Compute squared differences
    mse_per_sample = tf.reduce_mean(tf.square(bbox_true - bbox_pred), axis=-1)

    # Condition: Set MSE to 0 for specific class labels
    condition = tf.logical_or(tf.argmax(class_true, axis=-1) == 1, tf.argmax(class_true,
axis=-1) == 3)
    mse_per_sample = tf.where(condition, 0.0, mse_per_sample)

    # Return mean MSE across the batch
    return tf.reduce_mean(mse_per_sample) * (10/7) # -> *(10/7) zato jer 30% podataka cl
ass(NOPALM ili UNDEFINED) pa kvare statistiku za cca ovaj omjer

early_stopping = EarlyStopping(
    monitor='val_loss', # Monitor validation loss
    patience=10, # Stop after 10 epochs of no improvement
    restore_best_weights=True, # Restore weights from the epoch with the best validation
loss
    verbose=1 # Print message when stopping
)

### GLOBAL #####
# lr_scheduler = ReduceLROnPlateau(
#     monitor='val_loss', # You can change this to 'val_iou' or another metric
#     factor=2/3, # Reduce learning rate by half
#     patience=5, # Number of epochs to wait for improvement before reducing th
e LR
#     min_lr= 0.00015, #1e-7 # Minimum learning rate

```

```

# verbose=1
# )

# optimizer = Adam(learning_rate=0.001)
# # Glob avg pool: start lr from: 0.0015, stop lr: 0.00015, patience=5, factor=2/3, just one: 0.001
# # Flatten: start lr: 0.000032, stop lr: 0.000002, just_one: 0.000008
#####

### FLATTEN #####
lr_scheduler = ReduceLROnPlateau(
    monitor='val_loss', # You can change this to 'val_iou' or another metric
    factor=0.5, # Reduce learning rate by half
    patience=3, # Number of epochs to wait for improvement before reducing the
LR
    min_lr= 0.00000002, #1e-7 # Minimum learning rate
    verbose=1
)

optimizer = Adam(learning_rate=0.000052) #optimizer = Adam(learning_rate=0.000032)
# Glob avg pool: start lr from: 0.0015, stop lr: 0.00015, patience=5, factor=2/3, just on
e: 0.001
# Flatten: start lr: 0.000032, stop lr: 0.000002, just_one: 0.000008
#####

# BUILD MODEL
model = build_model((IMG_SIZE[0], IMG_SIZE[1], 1), isMOVM = include_MOVM)

model.compile(
    optimizer=optimizer,
    loss=total_loss, # Using custom loss
    metrics=[
        class_accuracy, # Tracks % of the total data that was correctly predic
ted
        iou_metric, # Track average IoU across the epoch
        mse_metric # Track Mean Squared Error (MSE)
    ]
)

# model.compile(
#     optimizer=optimizer,
#     #loss=total_loss, # Using custom loss
#     loss={
#         'class_output': custom_scc_loss,
#         'comb_output': total_loss,
#     },
#     loss_weights={
#         'class_output': 0.5,
#         'comb_output': 1.0,
#     },
#     # metrics=[
#     #     #BinaryAccuracy(name="class_accuracy"), # Track binary classification accur
acy
#     #     class_accuracy, # Tracks % of the total data that was correctly predic
ted
#     #     iou_metric, # Track average IoU across the epoch
#     #     mse_metric # Track Mean Squared Error (MSE)
#     # ]
#     metrics={
#         'class_output': [class_accuracy], # Metrics for the class_output (classificati
on)
#         'comb_output': [iou_metric, mse_metric], # Metrics for the bbox_output (boundi
ng box)
#     }
# )

# Train the model
history = model.fit(

```

```
[X_train[..., 0], X_train[..., 1]] if include_MOVM else X_train,
y_train,
validation_data=([X_test[..., 0], X_test[..., 1]] if include_MOVM else X_test, y_test),
# y_test also contains the same structure
epochs=150, #150, #50,
batch_size=1,
callbacks=[lr_scheduler, early_stopping]
)
```

Model: "functional\_7"

Layer (type)	Output Shape	Param #	Connected to
input_layer_7 (InputLayer)	(None, 240, 360, 1)	0	-
cast_7 (Cast)	(None, 240, 360, 1)	0	input_layer_7[0][0]
conv2d_42 (Conv2D)	(None, 120, 180, 64)	3,200	cast_7[0][0]
leaky_re_lu_42 (LeakyReLU)	(None, 120, 180, 64)	0	conv2d_42[0][0]
max_pooling2d_21 (MaxPooling2D)	(None, 60, 90, 64)	0	leaky_re_lu_42[0]
conv2d_43 (Conv2D)	(None, 60, 90, 128)	73,856	max_pooling2d_21[0][0]
leaky_re_lu_43 (LeakyReLU)	(None, 60, 90, 128)	0	conv2d_43[0][0]
max_pooling2d_22 (MaxPooling2D)	(None, 30, 45, 128)	0	leaky_re_lu_43[0]
conv2d_44 (Conv2D)	(None, 30, 45, 64)	8,256	max_pooling2d_22[0][0]
leaky_re_lu_44	(None, 30, 45, 64)	0	conv2d_44[0][0]

(LeakyReLU)			
conv2d_45 (Conv2D)	(None, 30, 45, 128)	73,856	leaky_re_lu_44[0]
leaky_re_lu_45	(None, 30, 45, 128)	0	conv2d_45[0][0]
(LeakyReLU)			
conv2d_46 (Conv2D)	(None, 30, 45, 128)	16,512	leaky_re_lu_45[0]
leaky_re_lu_46	(None, 30, 45, 128)	0	conv2d_46[0][0]
(LeakyReLU)			
conv2d_47 (Conv2D)	(None, 30, 45, 128)	147,584	leaky_re_lu_46[0]
leaky_re_lu_47	(None, 30, 45, 128)	0	conv2d_47[0][0]
(LeakyReLU)			
max_pooling2d_23	(None, 15, 22, 128)	0	leaky_re_lu_47[0]
(MaxPooling2D)			
flatten_7 (Flatten)	(None, 42240)	0	max_pooling2d_23[0][0]
dense_7 (Dense)	(None, 128)	5,406,848	flatten_7[0][0]
dense_8 (Dense)	(None, 64)	8,256	dense_7[0][0]
class_output (Dense)	(None, 4)	260	dense_8[0][0]
bbox_output (Dense)	(None, 4)	516	dense_7[0][0]
output (Concatenate)	(None, 8)	0	class_output[0][0], bbox_output[0][0]

**Total params: 5,739,144 (21.89 MB)**

**Trainable params: 5,739,144 (21.89 MB)**

**Non-trainable params: 0 (0.00 B)**

Epoch 1/150

1006/1006 16s 8ms/step - class\_accuracy: 0.3761 - iou\_metric: 0.0000e+00 - loss: 20627.0156 - mse\_metric: 7322.3569 - val\_class\_accuracy: 0.3611 - val\_iou\_metric: 0.0000e+00 - val\_loss: 8531.4385 - val\_mse\_metric: 3014.0027 - learning\_rate: 5.2000e-05

Epoch 2/150

1006/1006 13s 5ms/step - class\_accuracy: 0.4077 - iou\_metric: 0.0000e+00 - loss: 9413.2627 - mse\_metric: 3328.8225 - val\_class\_accuracy: 0.2616 - val\_iou\_metric: 0.0000e+00 - val\_loss: 8088.9614 - val\_mse\_metric: 2855.9639 - learning\_rate: 5.2000e-05

Epoch 3/150

1006/1006 11s 6ms/step - class\_accuracy: 0.4030 - iou\_metric: 0.0000e+00 - loss: 8968.0938 - mse\_metric: 3170.2717 - val\_class\_accuracy: 0.4514 - val\_iou\_metric: 0.0000e+00 - val\_loss: 7371.6133 - val\_mse\_metric: 2599.5134 - learning\_rate: 5.2000e-05

Epoch 4/150

1006/1006 9s 5ms/step - class\_accuracy: 0.4102 - iou\_metric: 0.0000e+00 - loss: 7607.3794 - mse\_metric: 2683.3328 - val\_class\_accuracy: 0.4144 - val\_iou\_metric: 0.0000e+00 - val\_loss: 6171.8535 - val\_mse\_metric: 2170.2578 - learning\_rate: 5.2000e-05

Epoch 5/150

1006/1006 6s 6ms/step - class\_accuracy: 0.4040 - iou\_metric: 0.0378 - loss: 5560.7803 - mse\_metric: 1974.3972 - val\_class\_accuracy: 0.3611 - val\_iou\_metric: 0.1504 - val\_loss: 2600.7688 - val\_mse\_metric: 989.4973 - learning\_rate: 5.2000e-05

Epoch 6/150

1006/1006 5s 5ms/step - class\_accuracy: 0.3889 - iou\_metric: 0.1571 - loss: 2707.0737 - mse\_metric: 1028.4290 - val\_class\_accuracy: 0.5069 - val\_iou\_metric: 0.1575 - val\_loss: 2440.6172 - val\_mse\_metric: 917.4199 - learning\_rate: 5.2000e-05

Epoch 7/150

1006/1006 10s 5ms/step - class\_accuracy: 0.4755 - iou\_metric: 0.1881 - loss: 2186.7153 - mse\_metric: 847.2201 - val\_class\_accuracy: 0.4653 - val\_iou\_metric: 0.2226 - val\_loss: 1740.8560 - val\_mse\_metric: 682.5159 - learning\_rate: 5.2000e-05

Epoch 8/150

1006/1006 11s 6ms/step - class\_accuracy: 0.5090 - iou\_metric: 0.2151 - loss: 1695.0743 - mse\_metric: 664.6793 - val\_class\_accuracy: 0.4190 - val\_iou\_metric: 0.2201 - val\_loss: 1666.9321 - val\_mse\_metric: 650.2409 - learning\_rate: 5.2000e-05

Epoch 9/150

1006/1006 9s 5ms/step - class\_accuracy: 0.5594 - iou\_metric: 0.2436 - loss: 1519.1278 - mse\_metric: 604.9079 - val\_class\_accuracy: 0.4907 - val\_iou\_metric: 0.2230 - val\_loss: 1752.8549 - val\_mse\_metric: 673.2213 - learning\_rate: 5.2000e-05

Epoch 10/150

1006/1006 7s 6ms/step - class\_accuracy: 0.5674 - iou\_metric: 0.2419 - loss: 1420.2056 - mse\_metric: 558.7230 - val\_class\_accuracy: 0.4630 - val\_iou\_metric: 0.2324 - val\_loss: 1634.2407 - val\_mse\_metric: 633.3495 - learning\_rate: 5.2000e-05

Epoch 11/150

1006/1006 5s 5ms/step - class\_accuracy: 0.6629 - iou\_metric: 0.2483 - loss: 1132.8005 - mse\_metric: 458.8701 - val\_class\_accuracy: 0.5995 - val\_iou\_metric: 0.2367 - val\_loss: 1427.0640 - val\_mse\_metric: 560.1278 - learning\_rate: 5.2000e-05

Epoch 12/150

1006/1006 10s 5ms/step - class\_accuracy: 0.6386 - iou\_metric: 0.3030 - loss: 978.1494 - mse\_metric: 407.9620 - val\_class\_accuracy: 0.6111 - val\_iou\_metric: 0.2764 - val\_loss: 1233.8810 - val\_mse\_metric: 482.3761 - learning\_rate: 5.2000e-05

Epoch 13/150

1006/1006 11s 6ms/step - class\_accuracy: 0.6952 - iou\_metric: 0.2971 - loss: 1000.9975 - mse\_metric: 411.6629 - val\_class\_accuracy: 0.6227 - val\_iou\_metric: 0.2817 - val\_loss: 1139.6327 - val\_mse\_metric: 450.3763 - learning\_rate: 5.2000e-05

Epoch 14/150

1006/1006 5s 5ms/step - class\_accuracy: 0.6811 - iou\_metric: 0.3159 - loss: 877.7830 - mse\_metric: 365.0973 - val\_class\_accuracy: 0.6088 - val\_iou\_metric: 0.2995 - val\_loss: 1057.1931 - val\_mse\_metric: 421.8191 - learning\_rate: 5.2000e-05

Epoch 15/150

1006/1006 5s 5ms/step - class\_accuracy: 0.7060 - iou\_metric: 0.3174 - loss: 726.0701 - mse\_metric: 306.4318 - val\_class\_accuracy: 0.6528 - val\_iou\_metric: 0.3174

2988 - val\_loss: 997.8541 - val\_mse\_metric: 403.2990 - learning\_rate: 5.2000e-05  
Epoch 16/150  
1006/1006 5s 5ms/step - class\_accuracy: 0.7245 - iou\_metric: 0.3517  
- loss: 682.0455 - mse\_metric: 293.0174 - val\_class\_accuracy: 0.6134 - val\_iou\_metric: 0.  
2696 - val\_loss: 1121.5245 - val\_mse\_metric: 441.1426 - learning\_rate: 5.2000e-05  
Epoch 17/150  
1006/1006 11s 6ms/step - class\_accuracy: 0.7146 - iou\_metric: 0.3131  
- loss: 616.7276 - mse\_metric: 261.6057 - val\_class\_accuracy: 0.6204 - val\_iou\_metric: 0.  
3088 - val\_loss: 888.9110 - val\_mse\_metric: 361.7423 - learning\_rate: 5.2000e-05  
Epoch 18/150  
1006/1006 5s 5ms/step - class\_accuracy: 0.7683 - iou\_metric: 0.3742  
- loss: 459.8280 - mse\_metric: 209.5753 - val\_class\_accuracy: 0.6319 - val\_iou\_metric: 0.  
2523 - val\_loss: 1087.8544 - val\_mse\_metric: 425.5473 - learning\_rate: 5.2000e-05  
Epoch 19/150  
1006/1006 5s 5ms/step - class\_accuracy: 0.7636 - iou\_metric: 0.3461  
- loss: 518.3761 - mse\_metric: 224.3706 - val\_class\_accuracy: 0.6736 - val\_iou\_metric: 0.  
2589 - val\_loss: 1036.5045 - val\_mse\_metric: 408.7679 - learning\_rate: 5.2000e-05  
Epoch 20/150  
1006/1006 6s 5ms/step - class\_accuracy: 0.7464 - iou\_metric: 0.3935  
- loss: 422.3975 - mse\_metric: 185.3858 - val\_class\_accuracy: 0.6620 - val\_iou\_metric: 0.  
3261 - val\_loss: 809.9001 - val\_mse\_metric: 328.3138 - learning\_rate: 5.2000e-05  
Epoch 21/150  
1006/1006 5s 5ms/step - class\_accuracy: 0.8008 - iou\_metric: 0.4319  
- loss: 344.5787 - mse\_metric: 162.2257 - val\_class\_accuracy: 0.6806 - val\_iou\_metric: 0.  
3458 - val\_loss: 749.7704 - val\_mse\_metric: 304.4341 - learning\_rate: 5.2000e-05  
Epoch 22/150  
1006/1006 5s 5ms/step - class\_accuracy: 0.8095 - iou\_metric: 0.4709  
- loss: 282.0826 - mse\_metric: 133.5281 - val\_class\_accuracy: 0.6250 - val\_iou\_metric: 0.  
3454 - val\_loss: 734.8336 - val\_mse\_metric: 297.0541 - learning\_rate: 5.2000e-05  
Epoch 23/150  
1006/1006 5s 5ms/step - class\_accuracy: 0.8027 - iou\_metric: 0.4607  
- loss: 285.9562 - mse\_metric: 138.0040 - val\_class\_accuracy: 0.6736 - val\_iou\_metric: 0.  
3555 - val\_loss: 702.4413 - val\_mse\_metric: 286.1792 - learning\_rate: 5.2000e-05  
Epoch 24/150  
1006/1006 5s 5ms/step - class\_accuracy: 0.8171 - iou\_metric: 0.4449  
- loss: 233.0018 - mse\_metric: 115.1150 - val\_class\_accuracy: 0.4352 - val\_iou\_metric: 0.  
3303 - val\_loss: 866.2593 - val\_mse\_metric: 330.8800 - learning\_rate: 5.2000e-05  
Epoch 25/150  
1006/1006 6s 6ms/step - class\_accuracy: 0.8318 - iou\_metric: 0.4920  
- loss: 231.5166 - mse\_metric: 112.5189 - val\_class\_accuracy: 0.6343 - val\_iou\_metric: 0.  
3317 - val\_loss: 753.8298 - val\_mse\_metric: 307.3468 - learning\_rate: 5.2000e-05  
Epoch 26/150  
1003/1006 0s 4ms/step - class\_accuracy: 0.8191 - iou\_metric: 0.4657  
- loss: 220.6829 - mse\_metric: 108.6716  
Epoch 26: ReduceLROnPlateau reducing learning rate to 2.5999999706982635e-05.  
1006/1006 9s 5ms/step - class\_accuracy: 0.8191 - iou\_metric: 0.4657  
- loss: 220.7547 - mse\_metric: 108.7026 - val\_class\_accuracy: 0.7060 - val\_iou\_metric: 0.  
3465 - val\_loss: 710.0319 - val\_mse\_metric: 292.0042 - learning\_rate: 5.2000e-05  
Epoch 27/150  
1006/1006 6s 6ms/step - class\_accuracy: 0.8766 - iou\_metric: 0.5321  
- loss: 132.4313 - mse\_metric: 71.3822 - val\_class\_accuracy: 0.6782 - val\_iou\_metric: 0.3  
695 - val\_loss: 669.4116 - val\_mse\_metric: 272.6316 - learning\_rate: 2.6000e-05  
Epoch 28/150  
1006/1006 5s 5ms/step - class\_accuracy: 0.8880 - iou\_metric: 0.5540  
- loss: 124.4042 - mse\_metric: 67.3281 - val\_class\_accuracy: 0.7014 - val\_iou\_metric: 0.3  
889 - val\_loss: 626.9677 - val\_mse\_metric: 257.9487 - learning\_rate: 2.6000e-05  
Epoch 29/150  
1006/1006 10s 5ms/step - class\_accuracy: 0.9073 - iou\_metric: 0.5923  
- loss: 97.7967 - mse\_metric: 56.2162 - val\_class\_accuracy: 0.7037 - val\_iou\_metric: 0.38  
47 - val\_loss: 610.2598 - val\_mse\_metric: 250.0912 - learning\_rate: 2.6000e-05  
Epoch 30/150  
1006/1006 11s 6ms/step - class\_accuracy: 0.8904 - iou\_metric: 0.5895  
- loss: 91.0143 - mse\_metric: 53.1640 - val\_class\_accuracy: 0.6968 - val\_iou\_metric: 0.38  
77 - val\_loss: 601.9846 - val\_mse\_metric: 247.6160 - learning\_rate: 2.6000e-05  
Epoch 31/150  
1006/1006 9s 5ms/step - class\_accuracy: 0.9198 - iou\_metric: 0.5735  
- loss: 96.8888 - mse\_metric: 56.1842 - val\_class\_accuracy: 0.6620 - val\_iou\_metric: 0.40  
44 - val\_loss: 575.3612 - val\_mse\_metric: 238.4911 - learning\_rate: 2.6000e-05  
Epoch 32/150  
1006/1006 6s 6ms/step - class\_accuracy: 0.9151 - iou\_metric: 0.5962  
- loss: 75.9031 - mse\_metric: 46.0760 - val\_class\_accuracy: 0.6481 - val\_iou\_metric: 0.39  
28 - val\_loss: 586.6464 - val\_mse\_metric: 242.4448 - learning\_rate: 2.6000e-05



Epoch 33/150  
1006/1006 5s 5ms/step - class\_accuracy: 0.9171 - iou\_metric: 0.5825  
- loss: 77.4031 - mse\_metric: 46.3078 - val\_class\_accuracy: 0.6875 - val\_iou\_metric: 0.2594 - val\_loss: 936.6805 - val\_mse\_metric: 370.9520 - learning\_rate: 2.6000e-05

Epoch 34/150  
1006/1006 5s 5ms/step - class\_accuracy: 0.9203 - iou\_metric: 0.5490  
- loss: 100.7263 - mse\_metric: 57.6677 - val\_class\_accuracy: 0.6782 - val\_iou\_metric: 0.4025 - val\_loss: 563.9953 - val\_mse\_metric: 233.9018 - learning\_rate: 2.6000e-05

Epoch 35/150  
1006/1006 6s 6ms/step - class\_accuracy: 0.9462 - iou\_metric: 0.6233  
- loss: 69.3015 - mse\_metric: 42.6301 - val\_class\_accuracy: 0.7014 - val\_iou\_metric: 0.3947 - val\_loss: 574.9741 - val\_mse\_metric: 237.9975 - learning\_rate: 2.6000e-05

Epoch 36/150  
1006/1006 10s 6ms/step - class\_accuracy: 0.9289 - iou\_metric: 0.5991  
- loss: 61.4953 - mse\_metric: 38.5578 - val\_class\_accuracy: 0.6435 - val\_iou\_metric: 0.3987 - val\_loss: 560.0239 - val\_mse\_metric: 233.3848 - learning\_rate: 2.6000e-05

Epoch 37/150  
1006/1006 9s 5ms/step - class\_accuracy: 0.9278 - iou\_metric: 0.6658  
- loss: 53.7378 - mse\_metric: 35.3933 - val\_class\_accuracy: 0.6921 - val\_iou\_metric: 0.4001 - val\_loss: 578.7515 - val\_mse\_metric: 238.8418 - learning\_rate: 2.6000e-05

Epoch 38/150  
1006/1006 6s 6ms/step - class\_accuracy: 0.9417 - iou\_metric: 0.6692  
- loss: 55.9984 - mse\_metric: 35.9875 - val\_class\_accuracy: 0.6991 - val\_iou\_metric: 0.4109 - val\_loss: 590.1279 - val\_mse\_metric: 241.1058 - learning\_rate: 2.6000e-05

Epoch 39/150  
996/1006 0s 4ms/step - class\_accuracy: 0.9350 - iou\_metric: 0.6405  
- loss: 51.8708 - mse\_metric: 34.0935

Epoch 39: ReduceLROnPlateau reducing learning rate to 1.2999999853491317e-05.

1006/1006 5s 5ms/step - class\_accuracy: 0.9349 - iou\_metric: 0.6405  
- loss: 51.8899 - mse\_metric: 34.1030 - val\_class\_accuracy: 0.6944 - val\_iou\_metric: 0.3976 - val\_loss: 584.3257 - val\_mse\_metric: 238.3082 - learning\_rate: 2.6000e-05

Epoch 40/150  
1006/1006 5s 5ms/step - class\_accuracy: 0.9412 - iou\_metric: 0.6855  
- loss: 33.2144 - mse\_metric: 23.1098 - val\_class\_accuracy: 0.6968 - val\_iou\_metric: 0.4160 - val\_loss: 547.9374 - val\_mse\_metric: 226.5429 - learning\_rate: 1.3000e-05

Epoch 41/150  
1006/1006 10s 5ms/step - class\_accuracy: 0.9724 - iou\_metric: 0.7323  
- loss: 31.4277 - mse\_metric: 22.3261 - val\_class\_accuracy: 0.6991 - val\_iou\_metric: 0.4174 - val\_loss: 538.9340 - val\_mse\_metric: 223.2766 - learning\_rate: 1.3000e-05

Epoch 42/150  
1006/1006 6s 6ms/step - class\_accuracy: 0.9465 - iou\_metric: 0.7254  
- loss: 26.0226 - mse\_metric: 19.0376 - val\_class\_accuracy: 0.6991 - val\_iou\_metric: 0.4116 - val\_loss: 553.3980 - val\_mse\_metric: 229.4954 - learning\_rate: 1.3000e-05

Epoch 43/150  
1006/1006 9s 5ms/step - class\_accuracy: 0.9631 - iou\_metric: 0.7087  
- loss: 26.7892 - mse\_metric: 19.7950 - val\_class\_accuracy: 0.6458 - val\_iou\_metric: 0.4156 - val\_loss: 546.3019 - val\_mse\_metric: 224.8112 - learning\_rate: 1.3000e-05

Epoch 44/150  
999/1006 0s 5ms/step - class\_accuracy: 0.9571 - iou\_metric: 0.7341  
- loss: 25.0281 - mse\_metric: 18.6892

Epoch 44: ReduceLROnPlateau reducing learning rate to 6.499999926745659e-06.

1006/1006 6s 5ms/step - class\_accuracy: 0.9571 - iou\_metric: 0.7339  
- loss: 25.0439 - mse\_metric: 18.6987 - val\_class\_accuracy: 0.6852 - val\_iou\_metric: 0.4194 - val\_loss: 547.8863 - val\_mse\_metric: 225.5279 - learning\_rate: 1.3000e-05

Epoch 45/150  
1006/1006 10s 5ms/step - class\_accuracy: 0.9712 - iou\_metric: 0.7331  
- loss: 18.4967 - mse\_metric: 14.5300 - val\_class\_accuracy: 0.6968 - val\_iou\_metric: 0.4240 - val\_loss: 535.9304 - val\_mse\_metric: 222.0641 - learning\_rate: 6.5000e-06

Epoch 46/150  
1006/1006 6s 6ms/step - class\_accuracy: 0.9803 - iou\_metric: 0.7550  
- loss: 17.4729 - mse\_metric: 13.6407 - val\_class\_accuracy: 0.6829 - val\_iou\_metric: 0.4280 - val\_loss: 538.5540 - val\_mse\_metric: 223.0716 - learning\_rate: 6.5000e-06

Epoch 47/150  
1006/1006 10s 6ms/step - class\_accuracy: 0.9865 - iou\_metric: 0.7665  
- loss: 18.5896 - mse\_metric: 15.0692 - val\_class\_accuracy: 0.6921 - val\_iou\_metric: 0.4244 - val\_loss: 532.3173 - val\_mse\_metric: 220.4612 - learning\_rate: 6.5000e-06

Epoch 48/150  
1006/1006 9s 5ms/step - class\_accuracy: 0.9825 - iou\_metric: 0.7651  
- loss: 16.0539 - mse\_metric: 13.0456 - val\_class\_accuracy: 0.7060 - val\_iou\_metric: 0.4274 - val\_loss: 533.2076 - val\_mse\_metric: 220.7447 - learning\_rate: 6.5000e-06

Epoch 49/150  
1006/1006 6s 6ms/step - class\_accuracy: 0.9853 - iou\_metric: 0.7377

1006/1006 5s 5ms/step - class\_accuracy: 0.9788 - iou\_metric: 0.7505  
- loss: 16.0919 - mse\_metric: 13.2210 - val\_class\_accuracy: 0.6991 - val\_iou\_metric: 0.4236 - val\_loss: 534.7366 - val\_mse\_metric: 221.6014 - learning\_rate: 6.5000e-06  
Epoch 50/150  
1006/1006 5s 5ms/step - class\_accuracy: 0.9788 - iou\_metric: 0.7505  
- loss: 16.0819 - mse\_metric: 13.1877 - val\_class\_accuracy: 0.7014 - val\_iou\_metric: 0.4230 - val\_loss: 530.6917 - val\_mse\_metric: 219.5492 - learning\_rate: 6.5000e-06  
Epoch 51/150  
1006/1006 5s 5ms/step - class\_accuracy: 0.9847 - iou\_metric: 0.7555  
- loss: 15.8945 - mse\_metric: 13.0344 - val\_class\_accuracy: 0.6875 - val\_iou\_metric: 0.4157 - val\_loss: 537.7115 - val\_mse\_metric: 222.7533 - learning\_rate: 6.5000e-06  
Epoch 52/150  
1006/1006 6s 5ms/step - class\_accuracy: 0.9866 - iou\_metric: 0.7649  
- loss: 15.3199 - mse\_metric: 12.5530 - val\_class\_accuracy: 0.6944 - val\_iou\_metric: 0.4245 - val\_loss: 533.1596 - val\_mse\_metric: 220.6167 - learning\_rate: 6.5000e-06  
Epoch 53/150  
1006/1006 10s 5ms/step - class\_accuracy: 0.9862 - iou\_metric: 0.7611  
- loss: 12.5356 - mse\_metric: 10.6450 - val\_class\_accuracy: 0.6991 - val\_iou\_metric: 0.4249 - val\_loss: 528.9001 - val\_mse\_metric: 218.8921 - learning\_rate: 6.5000e-06  
Epoch 54/150  
1006/1006 10s 5ms/step - class\_accuracy: 0.9807 - iou\_metric: 0.7597  
- loss: 14.4797 - mse\_metric: 11.5935 - val\_class\_accuracy: 0.6991 - val\_iou\_metric: 0.4195 - val\_loss: 537.5649 - val\_mse\_metric: 221.7848 - learning\_rate: 6.5000e-06  
Epoch 55/150  
1006/1006 6s 6ms/step - class\_accuracy: 0.9814 - iou\_metric: 0.7822  
- loss: 13.3445 - mse\_metric: 11.2500 - val\_class\_accuracy: 0.7060 - val\_iou\_metric: 0.4191 - val\_loss: 535.7968 - val\_mse\_metric: 221.8777 - learning\_rate: 6.5000e-06  
Epoch 56/150  
1002/1006 0s 4ms/step - class\_accuracy: 0.9892 - iou\_metric: 0.8056  
- loss: 13.4284 - mse\_metric: 11.4169  
Epoch 56: ReduceLROnPlateau reducing learning rate to 3.2499999633728294e-06.  
1006/1006 5s 5ms/step - class\_accuracy: 0.9891 - iou\_metric: 0.8054  
- loss: 13.4281 - mse\_metric: 11.4164 - val\_class\_accuracy: 0.7037 - val\_iou\_metric: 0.4207 - val\_loss: 534.8378 - val\_mse\_metric: 220.9633 - learning\_rate: 6.5000e-06  
Epoch 57/150  
1006/1006 6s 5ms/step - class\_accuracy: 0.9946 - iou\_metric: 0.7923  
- loss: 10.1104 - mse\_metric: 8.8584 - val\_class\_accuracy: 0.6991 - val\_iou\_metric: 0.4264 - val\_loss: 530.1531 - val\_mse\_metric: 219.4250 - learning\_rate: 3.2500e-06  
Epoch 58/150  
1006/1006 10s 5ms/step - class\_accuracy: 0.9914 - iou\_metric: 0.7819  
- loss: 11.0332 - mse\_metric: 9.5416 - val\_class\_accuracy: 0.7037 - val\_iou\_metric: 0.4238 - val\_loss: 530.3414 - val\_mse\_metric: 219.4827 - learning\_rate: 3.2500e-06  
Epoch 59/150  
1003/1006 0s 5ms/step - class\_accuracy: 0.9917 - iou\_metric: 0.7716  
- loss: 10.6909 - mse\_metric: 9.2749  
Epoch 59: ReduceLROnPlateau reducing learning rate to 1.6249999816864147e-06.  
1006/1006 6s 6ms/step - class\_accuracy: 0.9916 - iou\_metric: 0.7717  
- loss: 10.6916 - mse\_metric: 9.2756 - val\_class\_accuracy: 0.7037 - val\_iou\_metric: 0.4218 - val\_loss: 531.4850 - val\_mse\_metric: 219.6875 - learning\_rate: 3.2500e-06  
Epoch 60/150  
1006/1006 9s 5ms/step - class\_accuracy: 0.9928 - iou\_metric: 0.8150  
- loss: 9.1175 - mse\_metric: 8.2592 - val\_class\_accuracy: 0.7060 - val\_iou\_metric: 0.4268 - val\_loss: 528.0565 - val\_mse\_metric: 218.3153 - learning\_rate: 1.6250e-06  
Epoch 61/150  
1006/1006 6s 6ms/step - class\_accuracy: 0.9931 - iou\_metric: 0.7880  
- loss: 9.5410 - mse\_metric: 8.4967 - val\_class\_accuracy: 0.6968 - val\_iou\_metric: 0.4259 - val\_loss: 530.1737 - val\_mse\_metric: 219.2425 - learning\_rate: 1.6250e-06  
Epoch 62/150  
1006/1006 5s 5ms/step - class\_accuracy: 0.9952 - iou\_metric: 0.8022  
- loss: 10.1751 - mse\_metric: 8.9312 - val\_class\_accuracy: 0.6944 - val\_iou\_metric: 0.4263 - val\_loss: 527.3990 - val\_mse\_metric: 218.0893 - learning\_rate: 1.6250e-06  
Epoch 63/150  
1006/1006 5s 5ms/step - class\_accuracy: 0.9941 - iou\_metric: 0.8169  
- loss: 9.3246 - mse\_metric: 8.4248 - val\_class\_accuracy: 0.6944 - val\_iou\_metric: 0.4270 - val\_loss: 528.9360 - val\_mse\_metric: 218.6078 - learning\_rate: 1.6250e-06  
Epoch 64/150  
1006/1006 6s 5ms/step - class\_accuracy: 0.9905 - iou\_metric: 0.8108  
- loss: 8.6568 - mse\_metric: 7.9662 - val\_class\_accuracy: 0.6921 - val\_iou\_metric: 0.4286 - val\_loss: 530.2296 - val\_mse\_metric: 219.2297 - learning\_rate: 1.6250e-06  
Epoch 65/150  
1001/1006 0s 4ms/step - class\_accuracy: 0.9930 - iou\_metric: 0.7801  
- loss: 8.8411 - mse\_metric: 7.9210  
Epoch 65: ReduceLROnPlateau reducing learning rate to 8.124999908432073e-07.

```

Epoch 66: ReduceLROnPlateau reducing learning rate to 1.6250e-07.
1006/1006 ————— 10s 5ms/step - class_accuracy: 0.9930 - iou_metric: 0.7803
- loss: 8.8440 - mse_metric: 7.9233 - val_class_accuracy: 0.6944 - val_iou_metric: 0.4275
- val_loss: 527.9388 - val_mse_metric: 218.3549 - learning_rate: 1.6250e-07
Epoch 66/150
1006/1006 ————— 10s 5ms/step - class_accuracy: 0.9991 - iou_metric: 0.8393
- loss: 8.9136 - mse_metric: 7.9937 - val_class_accuracy: 0.6944 - val_iou_metric: 0.4272
- val_loss: 528.8261 - val_mse_metric: 218.3917 - learning_rate: 8.1250e-07
Epoch 67/150
1006/1006 ————— 10s 5ms/step - class_accuracy: 0.9932 - iou_metric: 0.8300
- loss: 8.2230 - mse_metric: 7.5977 - val_class_accuracy: 0.7014 - val_iou_metric: 0.4270
- val_loss: 528.0361 - val_mse_metric: 218.2209 - learning_rate: 8.1250e-07
Epoch 68/150
1002/1006 ————— 0s 4ms/step - class_accuracy: 0.9948 - iou_metric: 0.7873
- loss: 8.3555 - mse_metric: 7.6161
Epoch 68: ReduceLROnPlateau reducing learning rate to 4.0624999542160367e-07.
1006/1006 ————— 10s 5ms/step - class_accuracy: 0.9948 - iou_metric: 0.7874
- loss: 8.3570 - mse_metric: 7.6171 - val_class_accuracy: 0.6991 - val_iou_metric: 0.4270
- val_loss: 529.5509 - val_mse_metric: 218.6602 - learning_rate: 8.1250e-07
Epoch 69/150
1006/1006 ————— 5s 5ms/step - class_accuracy: 0.9954 - iou_metric: 0.8207
- loss: 8.9603 - mse_metric: 7.8578 - val_class_accuracy: 0.6968 - val_iou_metric: 0.4282
- val_loss: 528.0030 - val_mse_metric: 218.2516 - learning_rate: 4.0625e-07
Epoch 70/150
1006/1006 ————— 6s 6ms/step - class_accuracy: 0.9939 - iou_metric: 0.8138
- loss: 8.8107 - mse_metric: 7.9107 - val_class_accuracy: 0.6991 - val_iou_metric: 0.4275
- val_loss: 527.8857 - val_mse_metric: 218.2142 - learning_rate: 4.0625e-07
Epoch 71/150
1002/1006 ————— 0s 4ms/step - class_accuracy: 0.9919 - iou_metric: 0.8188
- loss: 7.2018 - mse_metric: 6.6233
Epoch 71: ReduceLROnPlateau reducing learning rate to 2.0312499771080184e-07.
1006/1006 ————— 5s 5ms/step - class_accuracy: 0.9919 - iou_metric: 0.8188
- loss: 7.2075 - mse_metric: 6.6281 - val_class_accuracy: 0.6968 - val_iou_metric: 0.4276
- val_loss: 528.0751 - val_mse_metric: 218.1803 - learning_rate: 4.0625e-07
Epoch 72/150
1006/1006 ————— 6s 5ms/step - class_accuracy: 0.9940 - iou_metric: 0.7890
- loss: 9.6437 - mse_metric: 8.3651 - val_class_accuracy: 0.6991 - val_iou_metric: 0.4279
- val_loss: 528.0644 - val_mse_metric: 218.2196 - learning_rate: 2.0312e-07
Epoch 72: early stopping
Restoring model weights from the end of the best epoch: 62.

```

In [16]:

```

# Vizualizacija procesa treniranja

# Get the training and validation metrics from the history object
train_iou = history.history['iou_metric']
val_iou = history.history['val_iou_metric']

# Plot the IoU over epochs
plt.figure(figsize=(12, 6))
plt.plot(train_iou, label='Train IoU')
plt.plot(val_iou, label='Validation IoU')
plt.title('IoU (Intersection over Union) over Epochs')
plt.xlabel('Epochs')
plt.ylabel('IoU')
plt.legend()
plt.show()

train_loss = history.history['loss']
val_loss = history.history['val_loss']

train_class_accuracy = history.history['class_accuracy']
val_class_accuracy = history.history['val_class_accuracy']

train_bbox_mse = history.history['mse_metric']
val_bbox_mse = history.history['val_mse_metric']

# Plot the loss over epochs
plt.figure(figsize=(12, 6))

```

```

# Loss plot
plt.subplot(1, 2, 1)
plt.plot(train_loss, label='Train Loss')
plt.plot(val_loss, label='Validation Loss')
plt.title('Loss over Epochs')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.ylim(0, 1500)
plt.legend()

# Accuracy plot
plt.subplot(1, 2, 2)
plt.plot(train_class_accuracy, label='Train Accuracy')
plt.plot(val_class_accuracy, label='Validation Accuracy')
plt.title('Classification Accuracy over Epochs')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()


plt.tight_layout()
plt.show()

# Optionally, you can also plot the bounding box MSE
plt.figure(figsize=(12, 6))
plt.plot(train_bbox_mse, label='Train BBox MSE')
plt.plot(val_bbox_mse, label='Validation BBox MSE')
plt.title('Bounding Box MSE over Epochs')
plt.xlabel('Epochs')
plt.ylabel('MSE')
plt.legend()
plt.ylim(0, 1000)
plt.show()


model.training = False

```


IoU (Intersection over Union) over Epochs

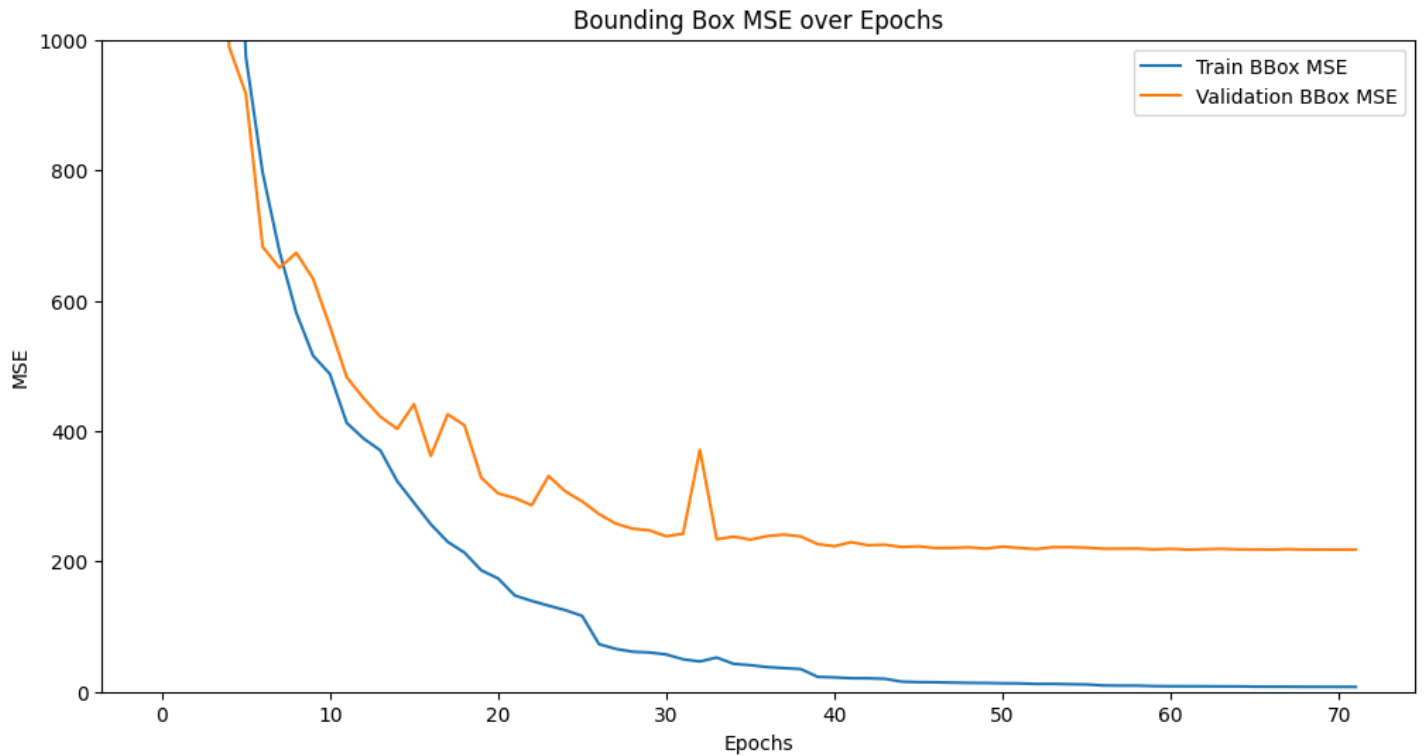



Loss over Epochs



Classification Accuracy over Epochs





In [51]:

```
# Praktično testiranje kako radi trenirani model, predikcije su prave predikcije (našeg)
# treniranog modela
# Parametar which_one služi da se odabere želiš li testirati na treniranim podacima, ili
# na dosad neviđenim test podacima

which_one = "test" # "train" or "test"

if which_one == "test":
    print("TEST")

    # Select 8 random indices from the filtered indices
    valid_indices = [
        i for i, label in enumerate(y_test[..., 0:4])
        if np.argmax(label[0:4]) == 0 or np.argmax(label[0:4]) == 2
    ]
    print("valid_indices num:", len(valid_indices))
    random_indices = random.sample(valid_indices, 8)

    # Loop through the random samples
    for index in random_indices:
        # Get the image, true bounding box, and true class label
        image = X_test[index, ..., 0] # Get grayscale image (remove channel dimension)
        true_bbox = y_test[index, 4:] # Get bounding box (without label)
        true_class = y_test[index, 0:4] # First element is the label (Palm or No Palm)

        # Get predicted bounding box and class from your model
        pred_output = model.predict(image[np.newaxis, ..., np.newaxis]) # Add batch and
        channel dimensions

        # Separate predicted class and bounding box from the model output
```

```

    pred_class = pred_output[0, 0:4] # First element is the predicted class (Palm or No Palm)
    pred_bbox = pred_output[0, 4:] # The remaining four elements are the predicted bounding box

    # Plot the image with both true and predicted bounding boxes
    plot_image_with_bbox(image, true_bbox, true_class, pred_bbox, pred_class, bbox_format=bbox_format)

elif which_one == "train":
    print("TRAIN")

    # Select 8 random indices from the filtered indices
    valid_indices = [
        i for i, label in enumerate(y_test[..., 0:4])
        if np.argmax(label[0:4]) == 0 or np.argmax(label[0:4]) == 2
    ]
    print("valid_indices num:", len(valid_indices))
    random_indices = random.sample(valid_indices, 8)

    # Loop through the random samples
    for index in random_indices:
        # Get the image, true bounding box, and true class label
        image = X_train[index, ..., 0] # Get grayscale image (remove channel dimension)
        true_bbox = y_train[index, 4:] # Get bounding box (without label)
        true_class = y_train[index, 0:4] # First element is the label (Palm or No Palm)

        # Get predicted bounding box and class from your model
        pred_output = model.predict(image[np.newaxis, ..., np.newaxis]) # Add batch and channel dimensions

        # Separate predicted class and bounding box from the model output
        pred_class = pred_output[0, 0:4] # First element is the predicted class (Palm or No Palm)
        pred_bbox = pred_output[0, 4:] # The remaining four elements are the predicted bounding box

        # Plot the image with both true and predicted bounding boxes
        plot_image_with_bbox(image, true_bbox, true_class, pred_bbox, pred_class, bbox_format=bbox_format)

```

```

TEST
valid_indices num: 289
1/1 ————— 0s 20ms/step
IoU: tf.Tensor(0.79221106, shape=(), dtype=float32)

```

True Class: [0. 0. 1. 0.] [PAUSE], Predicted Class: [5.035e-02 2.801e-06 9.497e-01 0.000e+00][PAUSE]



```


1/1 ————— 0s 19ms/step
IoU: tf.Tensor(0.70008194, shape=(), dtype=float32)

```

True Class: [1. 0. 0. 0.] [FLY], Predicted Class: [1.0e+00 0.0e+00 4.5e-06 0.0e+00][FLY]








1/1 ————— 0s 19ms/step  
 IoU: tf.Tensor(0.7675164, shape=(), dtype=float32)

True Class: [1. 0. 0. 0.] [FLY], Predicted Class: [1.00e+00 0.00e+00 5.96e-08 9.46e-05][FLY]



1/1 ————— 0s 19ms/step  
 IoU: tf.Tensor(0.6748983, shape=(), dtype=float32)


True Class: [0. 0. 1. 0.] [PAUSE], Predicted Class: [1.013e-06 1.994e-02 9.800e-01 5.901e-06][PAUSE]



1/1 ————— 0s 19ms/step  
 IoU: tf.Tensor(0.6881364, shape=(), dtype=float32)

True Class: [1. 0. 0. 0.] [FLY], Predicted Class: [9.990e-01 0.000e+00 5.960e-08 5.927e-04][FLY]






1/1 ————— 0s 18ms/step  
 IoU: tf.Tensor(0.8511985, shape=(), dtype=float32)

True Class: [0. 0. 1. 0.] [PAUSE], Predicted Class: [0.03906 0.004814 0.822 0.1342 ][[PAUSE]




1/1 ————— 0s 18ms/step  
 IoU: tf.Tensor(0.593128, shape=(), dtype=float32)

True Class: [1. 0. 0. 0.] [FLY], Predicted Class: [9.883e-01 0.000e+00 1.138e-05 1.169e-02][[FLY]




1/1 ————— 0s 23ms/step  
 IoU: tf.Tensor(0.23467761, shape=(), dtype=float32)

True Class: [1. 0. 0. 0.] [FLY], Predicted Class: [6.85e-06 4.77e-07 1.00e+00 0.00e+00][[PAUSE]







In [25]:

```
# Detaljna evaluacija IOU i MSE metrike (i TRAIN i TEST data)

def calculate_iou_mse_scores(dataset, model, class_filter, bbox_format_rn=None):
    """
    Calculate MSE scores for samples in the dataset where the class is in class_filter.

    Args:
    - dataset: Tuple (X, y) where X contains images and y contains labels and bounding boxes.
    - model: The trained model to make predictions.
    - class_filter: List of class labels to include in the MSE calculation (e.g., [0, 2])
    - bbox_format_rn: Bounding box format, e.g., 'center' or 'coco'. Default is None.

    Returns:
    - mse_scores: List of MSE scores for the filtered dataset.
    """
    X, y = dataset
    iou_just_0_2 = []
    mse_just_0_2 = []

    for i in range(len(X)):
        true_class = y[i, 0:4] # Get the class label
        true_class = np.argmax(true_class)
        if true_class in class_filter:

            # Get true bounding box
            true_bbox = y[i, 4:] # Bounding box (without label)

            # Predict bounding box and class
            image = X[i, ..., 0][np.newaxis, ..., np.newaxis] # Prepare image for prediction
            pred_output = model.predict(image, verbose=0)
            pred_bbox = pred_output[0, 4:] # Extract predicted bounding box

            # Calculate MSE (Mean Squared Error)
            mse = np.mean((true_bbox - pred_bbox) ** 2) # MSE between true and predicted bounding box
            if (mse > 1e6):
                print("pred output:", pred_output)
                print("mse0", mse)
                print()

            mse_just_0_2.append(float(mse))

            # Calculate IoU
            iou = 1 - iou_loss(true_bbox[np.newaxis, ...], pred_bbox[np.newaxis, ...], bbox_format_rn)
            iou_just_0_2.append(float(iou))

    return iou_just_0_2, mse_just_0_2

def plot_distributions(iou_scores, mse_scores):
```

```

# Calculate min and max for MSE x-axis limits
min_mse = min(mse_scores)
max_mse = max(mse_scores)

# Create subplots
fig, axes = plt.subplots(1, 2, figsize=(16, 6)) # 1 row, 2 columns

# IoU distribution plot
axes[0].hist(iou_scores, bins=20, range=(0, 1), color='blue', alpha=0.7, edgecolor='
black')
axes[0].set_title('Distribution of IoU Scores', fontsize=16)
axes[0].set_xlabel('IoU Score', fontsize=14)
axes[0].set_ylabel('Frequency', fontsize=14)
axes[0].grid(axis='y', linestyle='--', alpha=0.7)

# MSE distribution plot
axes[1].hist(mse_scores, bins=20, range=(min_mse, max_mse*0.3), color='blue', alpha=
0.7, edgecolor='black')
axes[1].set_title('Distribution of MSE Scores', fontsize=16)
axes[1].set_xlabel('MSE Score', fontsize=14)
axes[1].set_ylabel('Frequency', fontsize=14)
axes[1].grid(axis='y', linestyle='--', alpha=0.7)

# Adjust layout and show the plots
plt.tight_layout()
plt.show()

for which_one in ["train", "test"]:
    print("\n\n")
    print("-----", which_one, "-----")

    if which_one == "test":
        dataset = (X_test, y_test)
    elif which_one == "train":
        dataset = (X_train, y_train)

    class_filter = [0, 2] # Include only classes 0 and 2
    iou_just_0_2, mse_just_0_2 = calculate_iou_mse_scores(dataset, model, class_filter,
bbox_format_rn=bbox_format)

    # IOU
    avg_iou = np.mean(iou_just_0_2)
    std_iou = np.std(iou_just_0_2) # Standard deviation of IoU
    total_samples = len(iou_just_0_2)
    print("IoU:")
    print("Total analyzed samples:", total_samples)
    print(f"Average IoU for JUST classes {class_filter} for dataset {which_one}: {avg_iou
:.4f}")
    print(f"Standard Deviation of IoU for JUST classes {class_filter}: {std_iou:.4f}")

    # MSE
    avg_mse = np.mean(mse_just_0_2)
    std_mse = np.std(mse_just_0_2) # Standard deviation of MSE
    print("\nMSE:")
    print(f"Average MSE for JUST classes {class_filter} for dataset {which_one}: {avg_mse
:.4f}")
    print(f"Standard Deviation of MSE for JUST classes {class_filter}: {std_mse:.4f}")
    print()

# Plot IoU and MSE distributions side by side
plot_distributions(iou_just_0_2, mse_just_0_2)

```

----- train -----

IoU:

Total analyzed samples: 695


Average IoU for JUST classes [0, 2] for dataset train: 0.8166

Standard Deviation of IoU for JUST classes [0, 2]: 0.1108

MSE:

Average MSE for JUST classes [0, 2] for dataset train: 8.2508

Standard Deviation of MSE for JUST classes [0, 2]: 8.3508



----- test -----

IoU:

Total analyzed samples: 289


Average IoU for JUST classes [0, 2] for dataset test: 0.4460

Standard Deviation of IoU for JUST classes [0, 2]: 0.2803

MSE:

Average MSE for JUST classes [0, 2] for dataset test: 228.2007

Standard Deviation of MSE for JUST classes [0, 2]: 351.7076



In [52]:

```
import numpy as np
import matplotlib.pyplot as plt
from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay, classification_report, accuracy_score

def process_dataset(dataset, model):
    """
    Process the dataset to get true and predicted classes.
    Args:
        dataset: Tuple of (X, y), where X is the input data and y contains the labels.
```

```

        model: Trained model to predict outputs.
Returns:
    all_true_classes: Array of true class indices.
    all_pred_classes: Array of predicted class indices.
    """
X, y = dataset
all_true_classes = []
all_pred_classes = []

for i in range(len(X)):
    true_class = y[i, 0:4] # Get the class label
    true_class = np.argmax(true_class)
    all_true_classes.append(true_class)

    # Predict bounding box and class
    image = X[i, ..., 0][np.newaxis, ..., np.newaxis] # Prepare image for prediction

    pred_output = model.predict(image, verbose=0)
    pred_class = pred_output[0, 0:4] # Extract predicted class probabilities

    pred_class = np.argmax(pred_class)
    all_pred_classes.append(pred_class)

return np.array(all_true_classes), np.array(all_pred_classes)

def plot_confusion_matrix(y_true, y_pred, names_of_class, dataset_name="Dataset"):
    """
    Computes and plots the confusion matrix with class names.
    Args:
        y_true: True class labels.
        y_pred: Predicted class labels.
        names_of_class: List of class names.
        dataset_name: Name of the dataset for labeling the plot.
    """
    # Compute confusion matrix
    cm = confusion_matrix(y_true, y_pred)
    disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=names_of_class)

    # Plot the confusion matrix
    disp.plot(cmap=plt.cm.Blues, xticks_rotation=45)
    plt.title(f"Confusion Matrix for {dataset_name}")
    plt.show()

def print_metrics(y_true, y_pred, names_of_class, dataset_name="Dataset"):
    """
    Prints additional metrics like classification report and accuracy.
    Args:
        y_true: True class labels.
        y_pred: Predicted class labels.
        names_of_class: List of class names.
        dataset_name: Name of the dataset for labeling metrics.
    """
    print(f"----- Metrics for {dataset_name} -----")
    print(f"Accuracy: {accuracy_score(y_true, y_pred):.4f}")
    print("\nClassification Report:")
    print(classification_report(y_true, y_pred, target_names=names_of_class))


def process_and_plot(data, model, names_of_class, dataset_name="Dataset"):
    """
    Processes the dataset, plots the confusion matrix, and prints additional metrics.
    Args:
        data: Tuple of (X, y), where X is the input data and y contains the labels.
        model: Trained model to predict outputs.
        names_of_class: List of class names.
        dataset_name: Name of the dataset for labeling the outputs.
    """
    # Process the dataset to extract true and predicted labels
    y_true, y_pred = process_dataset(data, model)

    # Plot confusion matrix
    plot_confusion_matrix(y_true, y_pred, names_of_class, dataset_name)

```

```
# Print additional metrics
print_metrics(y_true, y_pred, names_of_class, dataset_name)


# Example usage
names_of_class = ["FLY", "NOPALM", "PAUSE", "UNDEFINED"]
process_and_plot((X_train, y_train), model, names_of_class, dataset_name="Train Dataset"
)
process_and_plot((X_test, y_test), model, names_of_class, dataset_name="Test Dataset")
```




----- Metrics for Train Dataset -----  
Accuracy: 0.9930

Classification Report:

	precision	recall	f1-score	support
FLY	0.99	1.00	1.00	527
NOPALM	1.00	1.00	1.00	137
PAUSE	0.99	0.99	0.99	168
UNDEFINED	0.99	0.97	0.98	174
accuracy			0.99	1006
macro avg	0.99	0.99	0.99	1006
weighted avg	0.99	0.99	0.99	1006





----- Metrics for Test Dataset -----

Accuracy: 0.6944

Classification Report:

	precision	recall	f1-score	support
FLY	0.80	0.83	0.82	215
NOPALM	0.82	0.78	0.80	63
PAUSE	0.56	0.62	0.59	74
UNDEFINED	0.39	0.33	0.36	80
accuracy			0.69	432
macro avg	0.64	0.64	0.64	432
weighted avg	0.69	0.69	0.69	432

In [ ]:

```
# calculate mAP?
```

In [ ]:

```
# Save the model, kasnije se može importat
model.save("coco_palm_detection_model_val_0.46_233.h5")
```

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save\_model(model)`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my\_model.keras')` or `keras.saving.save\_model(model, 'my\_model.keras')`.